

```
!pip install torch
```



```
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.me
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.me
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadat
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metada
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadat
Collecting nvidia-curand-cu12==10.3.5.147 (from torch)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.meta
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.meta
Collecting nvidia-cuspars-cu12==12.3.1.170 (from torch)
  Downloading nvidia_cuspars-cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.me
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.met
Requirement already satisfied: triton==3.1.0 in /usr/local/lib/python3.11/dist-pac
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-pac
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dis
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-p
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 M
    _____ 363.4/363.4 MB 2.9 MB/s eta 0:00:00
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.
    _____ 13.8/13.8 MB 78.3 MB/s eta 0:00:00
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.
    _____ 24.6/24.6 MB 61.3 MB/s eta 0:00:00
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (8
    _____ 883.7/883.7 kB 53.7 MB/s eta 0:00:00
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB
    _____ 664.8/664.8 MB 1.3 MB/s eta 0:00:00
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB
    _____ 211.5/211.5 MB 6.7 MB/s eta 0:00:00
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3
    _____ 56.3/56.3 MB 17.7 MB/s eta 0:00:00
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9
    _____ 127.9/127.9 MB 7.2 MB/s eta 0:00:00
  Downloading nvidia_cuspars-cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (207
    _____ 207.5/207.5 MB 6.1 MB/s eta 0:00:00
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1
    _____ 21.1/21.1 MB 95.1 MB/s eta 0:00:00
Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12, nvidia-c
  Attempting uninstall: nvidia-nvjitlink-cu12
    Found existing installation: nvidia-nvjitlink-cu12 12.5.82
      Uninstalling nvidia-nvjitlink-cu12 12.5.82:
```

```
Uninstalling nvidia-nvjitlink-cu12-12.5.82:  
Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82  
Attempting uninstall: nvidia-curand-cu12
```

```
!pip install torchinfo
```



```
Collecting torchinfo  
  Downloading torchinfo-1.8.0-py3-none-any.whl.metadata (21 kB)  
  Downloading torchinfo-1.8.0-py3-none-any.whl (23 kB)  
Installing collected packages: torchinfo  
Successfully installed torchinfo-1.8.0
```

```
# Creating U-Net architecture
```

```
# page 4 of the paper: (https://arxiv.org/pdf/1505.04597.pdf)
```

```
# Network Architecture
```

```
# The network architecture is illustrated in Figure 1. It consists of a contracting  
# path (left side) and an expansive path (right side). The contracting path follows  
# the typical architecture of a convolutional network. It consists of the repeated  
# application of two 3x3 convolutions (unpadded convolutions), each followed by  
# a rectified linear unit (ReLU) and a 2x2 max pooling operation with stride 2  
# for downsampling. At each downsampling step we double the number of feature  
# channels. Every step in the expansive path consists of an upsampling of the  
# feature map followed by a 2x2 convolution ("up-convolution") that halves the  
# number of feature channels, a concatenation with the correspondingly cropped  
# feature map from the contracting path, and two 3x3 convolutions, each followed by a Re  
# every convolution. At the final layer a 1x1 convolution is used to map each 64-  
# component feature vector to the desired number of classes. In total the network  
# has 23 convolutional layers.  
# To allow a seamless tiling of the output segmentation map (see Figure 2), it  
# is important to select the input tile size such that all 2x2 max-pooling operations  
# are applied to a layer with an even x- and y-size.
```

```
# The reduced output size within a single tile (e.g., 388x388 for a 572x572 input) ensur  
# avoiding incomplete or invalid segmentations near the borders.
```

```
#pytorch libraries
```

```
import torch  
import torch.nn as nn  
import torch.nn.functional as F #for ReLU  
import torchvision.transforms.functional as Trans  
from torchinfo import summary
```

```
class UNet(nn.Module):  
    def __init__(self, input_number, output_number):  
        super(UNet, self).__init__()  
        self.input_number = input_number  
        self.output_number = output_number
```

```

# Encoder
# input: 4d tensor: batch_size x input_number x 572x572 --> input number=1 for g
# assuming 572x572 image
self.conv1 = nn.Conv2d(self.input_number, 64, kernel_size=3, padding=0) # 64x572x572
self.conv2 = nn.Conv2d(64, 64, kernel_size=3, padding=0) # 64x568x568
self.maxPool1 = nn.MaxPool2d(kernel_size=2, stride=2) # 64x284x284

self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=0) # 128x282x282
self.conv4 = nn.Conv2d(128, 128, kernel_size=3, padding=0) # 128x280x280
self.maxPool2 = nn.MaxPool2d(kernel_size=2, stride=2) # 128x140x140

self.conv5 = nn.Conv2d(128, 256, kernel_size=3, padding=0) # 256x138x138
self.conv6 = nn.Conv2d(256, 256, kernel_size=3, padding=0) # 256x136x136
self.maxPool3 = nn.MaxPool2d(kernel_size=2, stride=2) # 256x68x68

self.conv7 = nn.Conv2d(256, 512, kernel_size=3, padding=0) # 512x66x66
self.conv8 = nn.Conv2d(512, 512, kernel_size=3, padding=0) # 512x64x64
self.maxPool4 = nn.MaxPool2d(kernel_size=2, stride=2) # 512x32x32

self.conv9 = nn.Conv2d(512, 1024, kernel_size=3, padding=0) # 1024x30x30
self.conv10 = nn.Conv2d(1024, 1024, kernel_size=3, padding=0) # 1024x28x28

# Decoder
# Upsampling by a factor of 2, --> stride=2, kernel_size=2
# 2x2 up convolution halves the feature channels

self.upconv1 = nn.ConvTranspose2d(1024, 512, kernel_size=2, stride=2) # 512x56x56
self.conv1b = nn.Conv2d(1024, 512, kernel_size=3, padding=0) # 512x54x54 other 5
self.conv2b = nn.Conv2d(512, 512, kernel_size=3, padding=0) # 512x52x52

self.upconv2 = nn.ConvTranspose2d(512, 256, 2, 2) # 512x104x104
self.conv3b = nn.Conv2d(512, 256, 3, padding=0) # 256x102x102
self.conv4b = nn.Conv2d(256, 256, 3, padding=0) # 256x100x100

self.upconv3 = nn.ConvTranspose2d(256, 128, 2, 2) # 256x200x200
self.conv5b = nn.Conv2d(256, 128, 3, padding=0) # 128x198x198
self.conv6b = nn.Conv2d(128, 128, 3, padding=0) # 128x196x196

self.upconv4 = nn.ConvTranspose2d(128, 64, 2, 2) # 128x392x392
self.conv7b = nn.Conv2d(128, 64, 3, padding=0) # 64x390x390
self.conv8b = nn.Conv2d(64, 64, 3, padding=0) # 64x388x388
self.final_conv = nn.Conv2d(64, self.output_number, kernel_size=1, padding=0) # 2

def cropConcat(self, encoder, decoder):
    # crops the encoder tensor and concatenate its with the decoder tensor
    _, _, H, W = decoder.shape
    cropped_enc = Trans.center_crop(encoder, [H, W]) # crops the encoder tensor in th
    return torch.cat((cropped_enc, decoder), dim=1) # concatenates at the feature di

def forward(self, x):
    # Encoder
    x = F.relu(self.conv1(x))
    x1 = F.relu(self.conv2(x))
    x = self.maxPool1(x1)

```

```

x = F.relu(self.conv3(x))
x2 = F.relu(self.conv4(x))
x = self.maxPool2(x2)

x = F.relu(self.conv5(x))
x3 = F.relu(self.conv6(x))
x = self.maxPool3(x3)

x = F.relu(self.conv7(x))
x4 = F.relu(self.conv8(x))
x = self.maxPool4(x4)

x = F.relu(self.conv9(x))
x = F.relu(self.conv10(x))

# Decoder
x = self.upconv1(x) # size 512x56x56

x = self.cropConcat(x4, x) # concatination1 size 1024x56x56
x = F.relu(self.conv1b(x))
x = F.relu(self.conv2b(x))
x = self.upconv2(x)

x = self.cropConcat(x3,x)
x = F.relu(self.conv3b(x))
x = F.relu(self.conv4b(x))
x = self.upconv3(x)

x = self.cropConcat(x2,x)
x = F.relu(self.conv5b(x))
x = F.relu(self.conv6b(x))
x = self.upconv4(x)

x = self.cropConcat(x1,x)
x = F.relu(self.conv7b(x))
x = F.relu(self.conv8b(x))
x = self.final_conv(x)
#x = F.softmax(x,1)
return x

```

```

model = UNet(1,1)
summary(model, input_size=(1, 1, 572, 572)) # Example input size

```

Layer (type:depth-idx)	Output Shape	Param #
UNet	[1, 1, 388, 388]	--
└─Conv2d: 1-1	[1, 64, 570, 570]	640
└─Conv2d: 1-2	[1, 64, 568, 568]	36,928
└─MaxPool2d: 1-3	[1, 64, 284, 284]	--
└─Conv2d: 1-4	[1, 128, 282, 282]	73,856
└─Conv2d: 1-5	[1, 128, 280, 280]	147,584

MaxPool2d: 1-6	[1, 128, 140, 140]	--
Conv2d: 1-7	[1, 256, 138, 138]	295,168
Conv2d: 1-8	[1, 256, 136, 136]	590,080
MaxPool2d: 1-9	[1, 256, 68, 68]	--
Conv2d: 1-10	[1, 512, 66, 66]	1,180,160
Conv2d: 1-11	[1, 512, 64, 64]	2,359,808
MaxPool2d: 1-12	[1, 512, 32, 32]	--
Conv2d: 1-13	[1, 1024, 30, 30]	4,719,616
Conv2d: 1-14	[1, 1024, 28, 28]	9,438,208
ConvTranspose2d: 1-15	[1, 512, 56, 56]	2,097,664
Conv2d: 1-16	[1, 512, 54, 54]	4,719,104
Conv2d: 1-17	[1, 512, 52, 52]	2,359,808
ConvTranspose2d: 1-18	[1, 256, 104, 104]	524,544
Conv2d: 1-19	[1, 256, 102, 102]	1,179,904
Conv2d: 1-20	[1, 256, 100, 100]	590,080
ConvTranspose2d: 1-21	[1, 128, 200, 200]	131,200
Conv2d: 1-22	[1, 128, 198, 198]	295,040
Conv2d: 1-23	[1, 128, 196, 196]	147,584
ConvTranspose2d: 1-24	[1, 64, 392, 392]	32,832
Conv2d: 1-25	[1, 64, 390, 390]	73,792
Conv2d: 1-26	[1, 64, 388, 388]	36,928
Conv2d: 1-27	[1, 1, 388, 388]	65

```

=====
Total params: 31,030,593
Trainable params: 31,030,593
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 167.45
=====

```

```

=====
Input size (MB): 1.31
Forward/backward pass size (MB): 1073.62
Params size (MB): 124.12
Estimated Total Size (MB): 1199.05
=====

```

```

# Script for visualizing and testing the U-net model
# Explanation of code in visualization_ReadMe.md

```

```

from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
import random
import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor, transforms
import torchvision.models as models
import torchvision.transforms.functional as Trans

```

```

# Count images in TIF file
def countImages(image):
    # Count images
    image_count = 0
    try:
        while True:
            image.seek(image_count)

```

```

        image_count +=1
    except EOFError:
        pass

    print(f"Number of pages: {image_count}\n")
    return image_count

# Display all Images of the dataset
def displayImages(image, title):
    plt.figure(figsize=(20,20))
    for i in range(countImages(image)):
        try:
            image.seek(i)
            plt.subplot(6,5, i+1)
            plt.imshow(image)
            plt.axis("off")
            plt.title(f"Page {i+1}")
        except EOFError:
            break

    plt.suptitle(title, fontsize=16)
    plt.tight_layout()
    plt.show(block=False)

# Prediction function for a single image
def predict_single(image, model):
    model.eval()
    image = image.unsqueeze(0).to(device) # Add batch dimension and move to device
    with torch.no_grad():
        pred = model(image)
    return pred

# function only for testing purposes of visualization
def evaluate_test(prediction, expected):
    probabilities = torch.sigmoid(prediction)
    probabilities = prediction
    predicted_segmentation = (probabilities > 0.5).int() #convert to binary mask (1,1

    _,H,W = predicted_segmentation.shape
    H=388
    W=388
    expected_segmentation = Trans.center_crop(expected, [H,W]) # crop the expected lab
    predicted_segmentation = Trans.center_crop(predicted_segmentation, [H,W]) # crop t

    diff = (expected_segmentation != predicted_segmentation).int()
    return predicted_segmentation, expected_segmentation, diff #return both binary ma

# correct evaluation function
def evaluate(prediction, expected):
    probabilities = torch.sigmoid(prediction)
    predicted_segmentation = (probabilities > 0.5).int() #convert to binary mask (1,1

    _,H,W = predicted_segmentation.shape

```

```

expected_segmentation = Trans.center_crop(expected, [H,W]) # crop the expected lab

diff = (expected_segmentation != predicted_segmentation).int() # difference mask
return predicted_segmentation, expected_segmentation, diff #return both binary ma

# selecting device and loading model
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
#model = torch.load('u_net.pth', weights_only=False)
#model = model.to(device)

# -----
# Displaying the complete train/test dataset
#Training images
train_image = Image.open("ISBI-2012-challenge/train-mirror.tif")
#displayImages(train_image, "Training images")
# Training labels
train_label = Image.open("ISBI-2012-challenge/train-labels.tif")
#displayImages(train_label, "Training labels")

# #Test images
test_image = Image.open("ISBI-2012-challenge/test-mirror.tif")
# displayImages(test_image, "Test images")
# Training labels
test_label = Image.open("ISBI-2012-challenge/test-labels.tif")
# displayImages(test_label, "Test labels")
#input("press enter for close")

#-----

# -----
# Visualization
# What we want to display: Image, Prediction; Groundtruth, Difference
# Image size = 512x512,
# Prediction size = 388x388
# select a random image of test set
randomID = random.randint(1, countImages(test_image))
test_image.seek(randomID)
test_label.seek(randomID)

plt.figure(figsize=(30,30))

# Test image
plt.subplot(2,2, 1)
plt.imshow(test_image)
plt.title(f"Test image")
plt.axis("off")

# Transform the test image/label to tensor
transform = transforms.Compose([transforms.ToTensor()]) # convert PIL image to (C,H,W)

```

```

image_tensor = transform(test_image)
label_tensor = transform(test_label)
#image_tensor = image_tensor.unsqueeze(0) # Shape (1,C,H,W)
#torch.set_printoptions(threshold=torch.inf)
print(f"Tensor size image_tensor: {image_tensor.shape}") # (1,1,512,512)
print(image_tensor)
print(f"Tensor size label tensor: {label_tensor.shape}") # (1,1,512,512)
#print(label_tensor)

# Run the U-net for one image
#prediction = predict_single(test_image, model)

# --> prediction_mask, label_mask, diff = evaluate(prediction, label_tensor)
prediction_mask, label_mask, diff = evaluate(image_tensor, label_tensor)
print(f"Tensor size prediction mask: {prediction_mask.shape}")
print(f"Tensor size label mask: {label_mask.shape}")
prediction_mask.squeeze_(0)
label_mask.squeeze_(0)
diff.squeeze_(0)
plt.subplot(2,2,2)
plt.imshow(prediction_mask, cmap='viridis')
plt.title(f"Prediction")
cbar = plt.colorbar()
cbar.set_label("0: membrane, 1: membrane")
plt.axis("off")
plt.subplot(2,2,3)
plt.imshow(label_mask, cmap='viridis')
cbar = plt.colorbar()
cbar.set_label("0: membrane, 1: cell")
plt.title(f"GT")
plt.axis("off")
plt.subplot(2,2,4)
plt.imshow(diff, cmap='viridis')
cbar = plt.colorbar()
cbar.set_label("0: correct, 1: wrong")
plt.title(f"Difference")
plt.axis("off")
# add colorbar

plt.show()

```

Number of pages: 30

```

Tensor size image_tensor: torch.Size([3, 572, 572])
tensor([[[[0.4314, 0.2745, 0.1608, ..., 0.5922, 0.5569, 0.4667],
          [0.4235, 0.3216, 0.2078, ..., 0.5059, 0.3922, 0.2627],
          [0.4863, 0.3529, 0.2196, ..., 0.5176, 0.4431, 0.3569],
          ...,
          [0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
          [0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
          [0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]],
        [[0.4314, 0.2745, 0.1608, ..., 0.5922, 0.5569, 0.4667],
          [0.4235, 0.3216, 0.2078, ..., 0.5059, 0.3922, 0.2627],
          [0.4863, 0.3529, 0.2196, ..., 0.5176, 0.4431, 0.3569],
          ...,
          [0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
          [0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
          [0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]]],
        [[0.4314, 0.2745, 0.1608, ..., 0.5922, 0.5569, 0.4667],
          [0.4235, 0.3216, 0.2078, ..., 0.5059, 0.3922, 0.2627],
          [0.4863, 0.3529, 0.2196, ..., 0.5176, 0.4431, 0.3569],
          ...,
          [0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
          [0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
          [0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]]],
        [[0.4314, 0.2745, 0.1608, ..., 0.5922, 0.5569, 0.4667],
          [0.4235, 0.3216, 0.2078, ..., 0.5059, 0.3922, 0.2627],
          [0.4863, 0.3529, 0.2196, ..., 0.5176, 0.4431, 0.3569],
          ...,
          [0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
          [0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
          [0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]]]])

```



```

...,
[0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
[0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
[0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]],

[[0.4314, 0.2745, 0.1608, ..., 0.5922, 0.5569, 0.4667],
 [0.4235, 0.3216, 0.2078, ..., 0.5059, 0.3922, 0.2627],
 [0.4863, 0.3529, 0.2196, ..., 0.5176, 0.4431, 0.3569],
 ...,
 [0.1843, 0.1529, 0.2078, ..., 0.4314, 0.4157, 0.4314],
 [0.1804, 0.1294, 0.2549, ..., 0.5882, 0.5020, 0.5490],
 [0.1608, 0.0745, 0.2471, ..., 0.6706, 0.6039, 0.5961]]])
Tensor size label tensor: torch.Size([1, 512, 512])
Tensor size prediction mask: torch.Size([3, 572, 572])
Tensor size label mask: torch.Size([1, 572, 572])

```

```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-4-2bfce6c38239> in <cell line: 0>()
    149 diff.squeeze_(0)
    150 plt.subplot(2,2,2)
--> 151 plt.imshow(prediction_mask, cmap='viridis')
    152 plt.title(f"Prediction")
    153 cbar = plt.colorbar()

```



4 frames

```

/usr/local/lib/python3.11/dist-packages/matplotlib/image.py in
_normalize_image_array(A)
    641         A = A.squeeze(-1) # If just (M, N, 1), assume scalar and
apply colormap.
    642         if not (A.ndim == 2 or A.ndim == 3 and A.shape[-1] in [3, 4]):
--> 643             raise TypeError(f"Invalid shape {A.shape} for image data")
    644         if A.ndim == 3:
    645             # If the input data has values outside the valid range (after

```

**TypeError:** Invalid shape (3, 572, 572) for image data



Next steps: [Explain error](#)

```
import torch
from torch import optim, nn
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms
import torchvision.transforms.functional as Trans
from PIL import Image
import matplotlib.pyplot as plt

torch.cuda.empty_cache()

#Loading data
#https://pytorch.org/tutorials/beginner/data_loading_tutorial.html
class Dataset(Dataset):
    def __init__(self, image_path, label_path, transform=None):
        self.images = Image.open(image_path)
        self.labels = Image.open(label_path)
        self.transform = transform

    def __len__(self):
        return self.images.n_frames

    def __getitem__(self, idx):
        #find specific frame
        self.images.seek(idx)
        self.labels.seek(idx)
        #grayscale conversion (if necessary)
        image = self.images.convert("L")
        label = self.labels.convert("L")
        if self.transform:
            image = self.transform(image)
            label = self.transform(label)
        return image, label

transform = transforms.Compose([transforms.ToTensor()])

#original
#train_dataset = Dataset("ISBI-2012-challenge/train-volume.tif", "ISBI-2012-challenge/tr
#test_dataset = Dataset("ISBI-2012-challenge/test-volume.tif", "ISBI-2012-challenge/test

#mirrored
train_dataset = Dataset("ISBI-2012-challenge/train-mirror.tif", "ISBI-2012-challenge/tr
test_dataset = Dataset("ISBI-2012-challenge/test-mirror.tif", "ISBI-2012-challenge/test-

train_loader = DataLoader(train_dataset, batch_size=1, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=1, shuffle=False)
```

```

# print(train_dataset.images.n_frames)

# print(train_dataset.shape)

# Training
# https://pytorch.org/tutorials/beginner/introyt/trainingyt.html
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = UNet(1,1).to(device)
optimizer = optim.Adam(model.parameters(), 0.001)
# optimizer = optim.SGD(model.parameters(), 0.001, momentum=0.99)
# criterion = nn.CrossEntropyLoss() # for some reason doesn't work
criterion = nn.BCEWithLogitsLoss()

def train(model, dataloader, criterion, optimizer, nrOfEpochs):
    model.train()
    for i in range(nrOfEpochs):
        running_loss = 0.0
        avgAcc=0.0
        avgPrec=0.0
        avgRec=0.0
        avgF1=0.0
        for images, labels in dataloader:
            images, labels = images.to(device), labels.to(device);
            # print(images.shape)
            optimizer.zero_grad()
            outputs=model(images)
            # outputs = (outputs > 0.5).float()
            labels = Trans.center_crop(labels, [388,388])
            # assert outputs.shape == labels.shape, f"Shape mismatch: {outputs.shape} vs {label
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
        if(i==nrOfEpochs-1):
            for j in range(dataloader.batch_size):
                outputs = (outputs > 0.5).int()
                labels = labels.int()
                TP = ((labels[j])*(outputs[j])).sum()
                TN = ((1-labels[j])*(1-outputs[j])).sum()
                FP = ((1-labels[j])*(outputs[j])).sum()
                FN = ((labels[j])*(1-outputs[j])).sum()

            accuracy = (TP+TN)/(TP+TN+FP+FN)
            precision = TP/(TP+FP)
            recall = TP/(TP+FN)
            f1 = 2*(precision*recall)/(precision+recall)
            ...

            print(f"Accuracy:{accuracy}")
            print(f"Precision:{precision}")
            print(f"Recall:{recall}")
            print(f"F1 Score:{f1}\n")
            ...

```

```

        avgAcc+=accuracy.cpu().item()
        avgPrec+=precision.cpu().item()
        avgRec+=recall.cpu().item()
        avgF1+=f1.cpu().item()

    ...
#for visualization of the training data
for j in range(dataloader.batch_size):
    plt.figure()
    plt.subplot(2,2,1)
    plt.imshow(images[j].cpu().numpy().squeeze(), cmap='viridis')
    plt.subplot(2,2,2)
    plt.imshow(labels[j].cpu().numpy().squeeze(), cmap='viridis')
    ...

train_loss = running_loss/len(dataloader)
print(f"Epoch {i+1}/{nrOfEpochs}\nLoss:{train_loss}")
return train_loss, avgAcc/len(dataloader), avgPrec/len(dataloader), avgRec/len(dataloader)

```

```

#Evaluation
def test(model, dataloader, criterion):
    model.eval()
    running_loss = 0.0
    avgAcc=0.0
    avgPrec=0.0
    avgRec=0.0
    avgF1=0.0
    plotImage=0
    with torch.no_grad():
        for images, labels in dataloader:
            images, labels = images.to(device), labels.to(device);
            outputs = model(images)
            labels = Trans.center_crop(labels, [388,388])
            loss = criterion(outputs,labels)
            running_loss+=loss.item()
            outputs = (outputs > 0.5).int()
            #visualisation and evaluation
            for j in range(dataloader.batch_size):
                if(plotImage==6): #for picking single image to plot
                    plt.figure()
                    plt.subplot(2,2,1)
                    plt.title(f"Image")
                    plt.imshow(images[j].cpu().numpy().squeeze(), cmap='viridis')
                    plt.subplot(2,2,2)
                    plt.title(f"Ground Truth")
                    plt.imshow(labels[j].cpu().numpy().squeeze(), cmap='viridis')
                    plt.subplot(2,2,3)
                    plt.title(f"Prediction")
                    #outputs[j]=(outputs[j]>0.5).int()
                    plt.imshow(outputs[j].cpu().detach().numpy().squeeze(), cmap='viridis')
                    plt.subplot(2,2,4)
                    plt.title(f"Difference")
                    diff=(outputs[j]!=labels[j]).int()
                    plt.imshow(diff.cpu().detach().numpy().squeeze(),cmap='viridis')

```

```

plt.tight_layout()
plt.show()

#Output image to binary values
#outputs[j] = (outputs > 0.5).float()
labels = labels.int()
TP = ((labels[j])*(outputs[j])).sum()
TN = ((1-labels[j])*(1-outputs[j])).sum()
FP = ((1-labels[j])*(outputs[j])).sum()
FN = ((labels[j])*(1-outputs[j])).sum()

accuracy = (TP+TN)/(TP+TN+FP+FN)
precision = TP/(TP+FP)
recall = TP/(TP+FN)
f1 = 2*(precision*recall)/(precision+recall)
'''

print(f"Accuracy:{accuracy}")
print(f"Precision:{precision}")
print(f"Recall:{recall}")
print(f"F1 Score:{f1}\n")
'''

avgAcc+=accuracy.cpu().item()
avgPrec+=precision.cpu().item()
avgRec+=recall.cpu().item()
avgF1+=f1.cpu().item()

plotImage+=1

print(f"Avg Accuracy:{avgAcc/len(dataloader)}")
print(f"Avg Precision:{avgPrec/len(dataloader)}")
print(f"Avg Recall:{avgRec/len(dataloader)}")
print(f"Avg F1 Score:{avgF1/len(dataloader)}\n")

test_loss = running_loss/len(dataloader)
print(f"Test loss:{test_loss}")
return test_loss, avgAcc/len(dataloader), avgPrec/len(dataloader), avgRec/len(dataloac

#10x number of epochs
nrEpx10 = 20

x = [0]*nrEpx10
trainLoss = [0]*nrEpx10
trainAcc = [0]*nrEpx10
trainPrec = [0]*nrEpx10
trainRec = [0]*nrEpx10
trainF1 = [0]*nrEpx10
testLoss = [0]*nrEpx10
testAcc = [0]*nrEpx10
testPrec = [0]*nrEpx10
testRec = [0]*nrEpx10
testF1 = [0]*nrEpx10

#torch.load('unet_50ep_lr0001_mirr.pth', map_location=torch.device('cpu'))

```

```

#Running code:
for i in range(nrEpx10):
    print(f"\nLoop {i+1}\n")
    x[i] = 10*(i+1)
    trainLoss[i], trainAcc[i], trainPrec[i], trainRec[i], trainF1[i] = train(model, train_
    testLoss[i], testAcc[i], testPrec[i], testRec[i], testF1[i] = test(model, test_loader,
    #update frame

#Plot acc, prec, rec, f1
fig = plt.figure()
#plt.title('Metrics over Epochs')
plt.subplot(2,3,1)
#plt.title(f"Loss")
plt.plot(x, trainLoss, "-o", label='Train Loss')
plt.plot(x, testLoss, "-o", label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.subplot(2,3,2)
#plt.title(f"Accuracy")
plt.plot(x, trainAcc, "-o", label='Train Accuracy')
plt.plot(x, testAcc, "-o", label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.subplot(2,3,3)
#plt.title(f"Precision")
plt.plot(x, trainPrec, "-o", label='Train Precision')
plt.plot(x, testPrec, "-o", label='Test Precision')
plt.xlabel('Epochs')
plt.ylabel('Precision')
plt.subplot(2,3,4)
#plt.title(f"Recall")
plt.plot(x, trainRec, "-o", label='Train Recall')
plt.plot(x, testRec, "-o", label='Test Recall')
plt.xlabel('Epochs')
plt.ylabel('Recall')
plt.subplot(2,3,5)
#plt.title(f"F1")
#plt.plot(x, trainF1, "-o", label='Train F1 Score')
#plt.plot(x, testF1, "-o", label='Test F1 Score')
plt.plot(x, trainF1, "-o", label='Train')
plt.plot(x, testF1, "-o", label='Test')
plt.xlabel('Epochs')
plt.ylabel('F1')
#plt.legend()
handles, labels = plt.gca().get_legend_handles_labels()
fig.legend(handles, labels, loc='lower right')
plt.tight_layout()
plt.show()

#Save model - change name to: unet _ nr of epochs _ learning rate 0.xxx as lrxxx _ which
#torch.save(model.state_dict(), 'unet_100ep_lr00001_mirr.pth')

```

Loop 1

Epoch 1/10

Loss:0.6694238523642222

Epoch 2/10

Loss:0.591159118215243

Epoch 3/10

Loss:0.528099466363589

Epoch 4/10

Loss:0.525952356060346

Epoch 5/10

Loss:0.5217139720916748

Epoch 6/10

Loss:0.5194161375363667

Epoch 7/10

Loss:0.489606378475825

Epoch 8/10

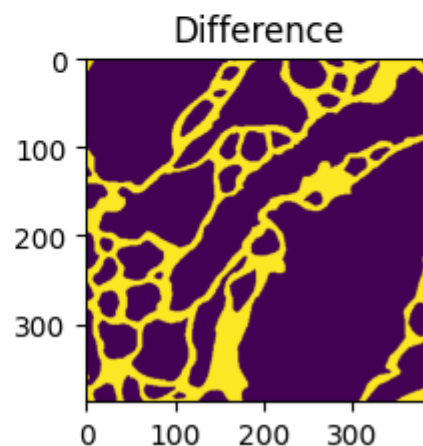

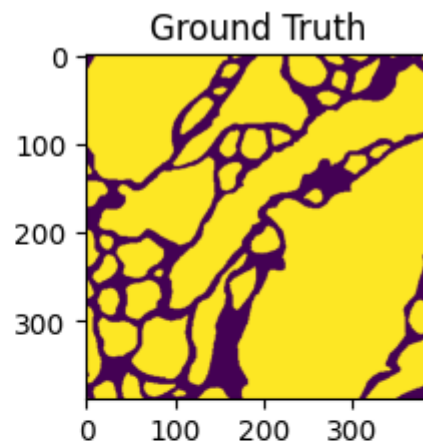

Loss:0.47040367821852364

Epoch 9/10

Loss:0.4394596715768178

Epoch 10/10

Loss:0.40733722150325774



Avg Accuracy:0.7215602477391561

Avg Precision:0.7215602477391561

Avg Recall:1.0

Avg F1 Score:0.8374458173910777

Test loss:0.5187153120835623

Loop 2

Epoch 1/10

Loss:0.3851223429044088

Epoch 2/10

Loss:0.3774399360020955

Epoch 3/10

Loss:0.38439785738786064

Epoch 4/10

Loss:0.3774010419845581

Epoch 5/10

Loss:0.35593909124533335

Epoch 6/10

Loss:0.34559838275114696

Epoch 7/10

Loss:0.36820192337036134

Epoch 8/10




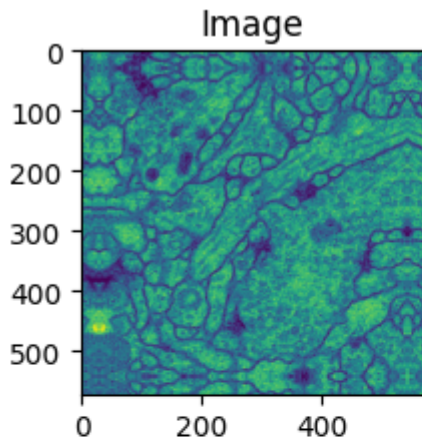
Loss:0.3587419251600901

Epoch 9/10

Loss:0.37042352159818015

Epoch 10/10

Loss:0.3790782610575358



Avg Accuracy:0.776786188284556

Avg Precision:0.9115005950133006

Avg Recall:0.7718825697898865

Avg F1 Score:0.8300122777620952

Test loss:0.42017103532950084

Loop 3


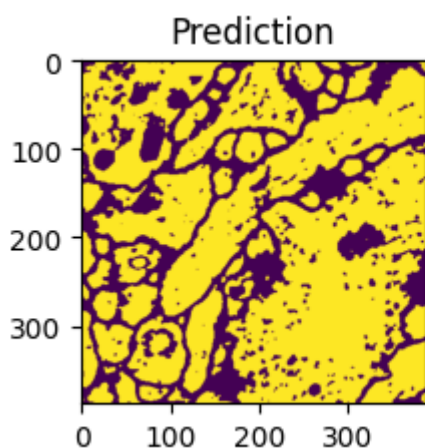


Epoch 1/10



```

Loss:0.36321892539660133
Epoch 2/10
Loss:0.3592783729235331
Epoch 3/10
Loss:0.35193413893381753
Epoch 4/10
Loss:0.3519824892282486
Epoch 5/10
Loss:0.3519705325365067
Epoch 6/10
Loss:0.3483248939116796
Epoch 7/10
Loss:0.3453931788603465
Epoch 8/10
Loss:0.3425674001375834
Epoch 9/10
Loss:0.3407235950231552
Epoch 10/10
Loss:0.34316649635632834

```



```

Avg Accuracy:0.7986314098040262
Avg Precision:0.8514258642991384
Avg Recall:0.8803501745065053
Avg F1 Score:0.8622707684834798

```

```
Test loss:0.4586367626984914
```





```
Loop 4
```

```

Epoch 1/10
Loss:0.3406229704618454
Epoch 2/10
Loss:0.3457350704076284

```

```
Loss:0.3457259794076284
Epoch 3/10
Loss:0.33847936590512595
Epoch 4/10
Loss:0.3443219393491745
Epoch 5/10
Loss:0.34111566146214806
Epoch 6/10
Loss:0.3407439102729162
Epoch 7/10
Loss:0.345645668109258
Epoch 8/10
Loss:0.34453535278638203
Epoch 9/10
Loss:0.34273766775925957
Epoch 10/10
Loss:0.33886585334936775
```



```
Avg Accuracy:0.8007448554039002
Avg Precision:0.8723234335581461
Avg Recall:0.8552560230096181
Avg F1 Score:0.859752082824707
```

```
Test loss:0.4245744655529658
```

```
Loop 5
```

```
Epoch 1/10
Loss:0.3403441200653712
Epoch 2/10
Loss:0.34240088164806365
Epoch 3/10
Loss:0.33833454648653666
```



































Start coding or [generate](#) with AI.